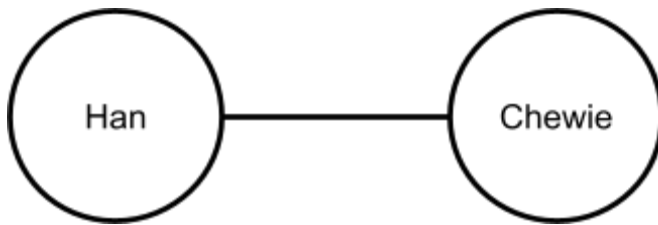


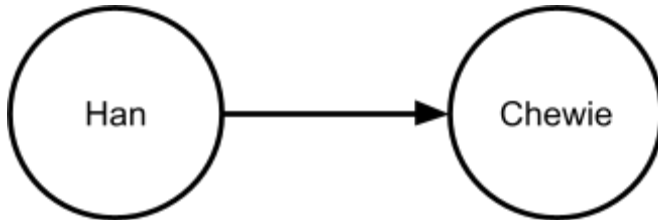
Graphs

- A **graph** is a structure representing a collection of objects or states, where some pairs of those objects are related or connected in some way.
- Graphs are used in lots and lots of places in computer science:
 - Social networks like Facebook or Twitter
 - Computer graphics
 - Machine learning
 - Computer vision
 - Logistics and optimization
 - Computer networking
- A graph is composed of **vertices** (or **nodes** or **points**) and **edges** (or **arcs** or **lines**).
- Vertices represent objects, states (i.e. conditions or configurations), locations, etc.
 - These form a set where each vertex is unique (i.e. no two vertices represent the same object/state): $V = \{v_1, v_2, v_3, \dots, v_n\}$
- Edges represent relationships or connections between vertices.
 - These are represented as vertex pairs: $E = \{(v_i, v_j), \dots\}$
 - Edges can be **directed** or **undirected**.
 - If there is an edge between v_i and v_j , then v_i and v_j are said to be **adjacent** (or they are **neighbors**).
 - Edges can be **weighted** or **unweighted**.
- An undirected edge is like a friend relationship in Facebook, e.g. if Han and Chewie are friends, there would be an undirected edge

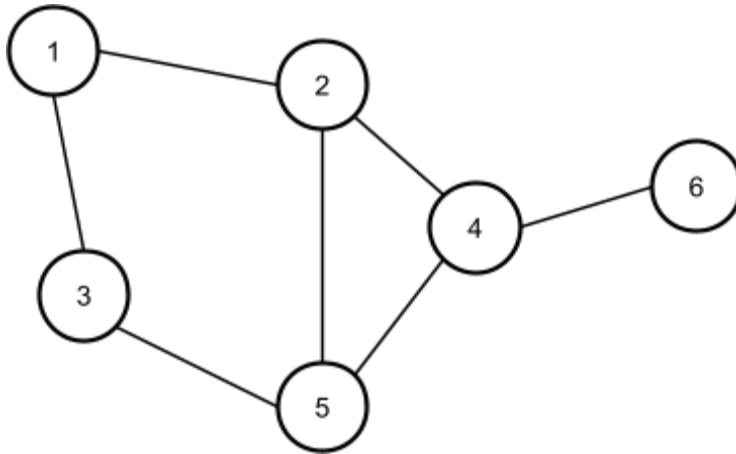
between them in the Facebook graph:



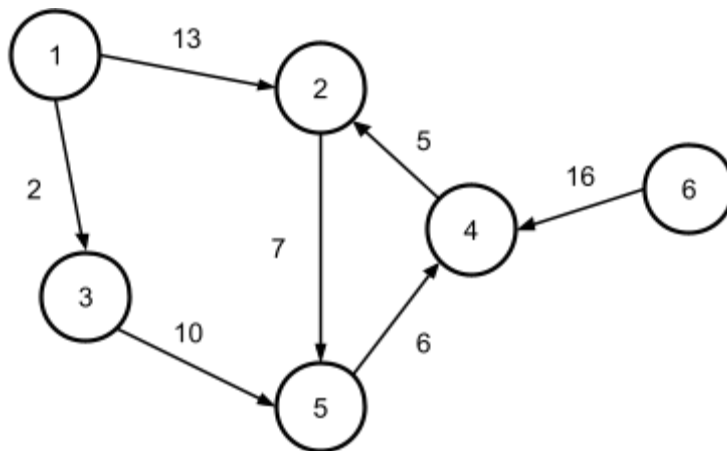
- A directed edge is like a follows relationship in Twitter, e.g. if Han follows Chewie, there would be a directed edge between them in the Twitter graph:



- Here, we say the edge is directed **from** Han **to** Chewie.
 - We can also say that Han is the **head** of this edge and that Chewie is its **tail**.
 - We can also say that Chewie is a **direct successor** of Han and that Han is a **direct predecessor** of Chewie.
 - We can also say that Chewie is **reachable** from Han.
- Here's an example of a small graph with 6 vertices and 7 undirected, unweighted edges:



- Here's an example of a similar graph with directed, weighted edges:

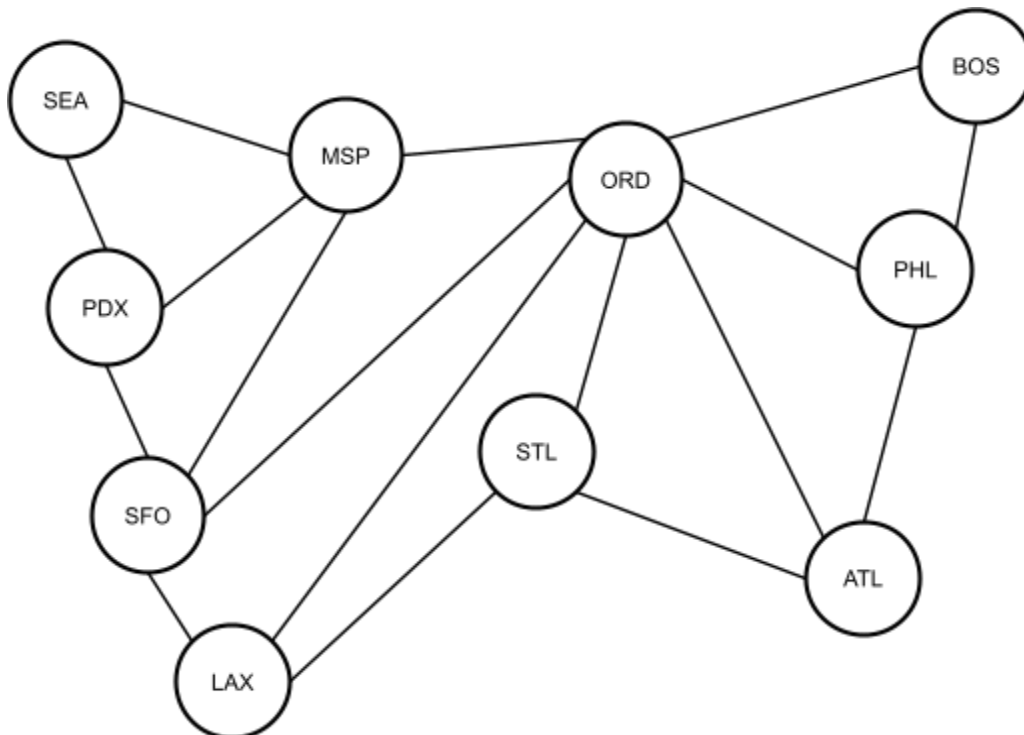


- Graphs represent general relationships between objects.
 - A node may have connections to any number of other nodes.
 - There can be multiple paths (or no path) from one node to another.
 - There can be cycles (loops) in the graph, where there is a path from one node back to itself.
- Trees are a special, more restricted subclass of graphs.
- There are lots of different kinds of questions we might want to ask about a graph:

- Is X in the graph?
- Is Y reachable from X?
- What nodes are reachable from X?
- Are X and Y adjacent?
- What's the shortest path from X to Y?
- How many edges between A and Y?

Representing graphs

- There are two main ways to represent a graph in practice:
 - An **adjacency list**, in which each vertex stores a list of its adjacent vertices.
 - An **adjacency matrix**, which is a two dimensional matrix whose rows and columns represent vertices. If there is an edge between v_i and v_j , the value at location (i, j) in the matrix will be non-zero.
- Let's consider this graph, where flights between US airports are represented, as an example:



- As an adjacency list, this graph would look like this:

```

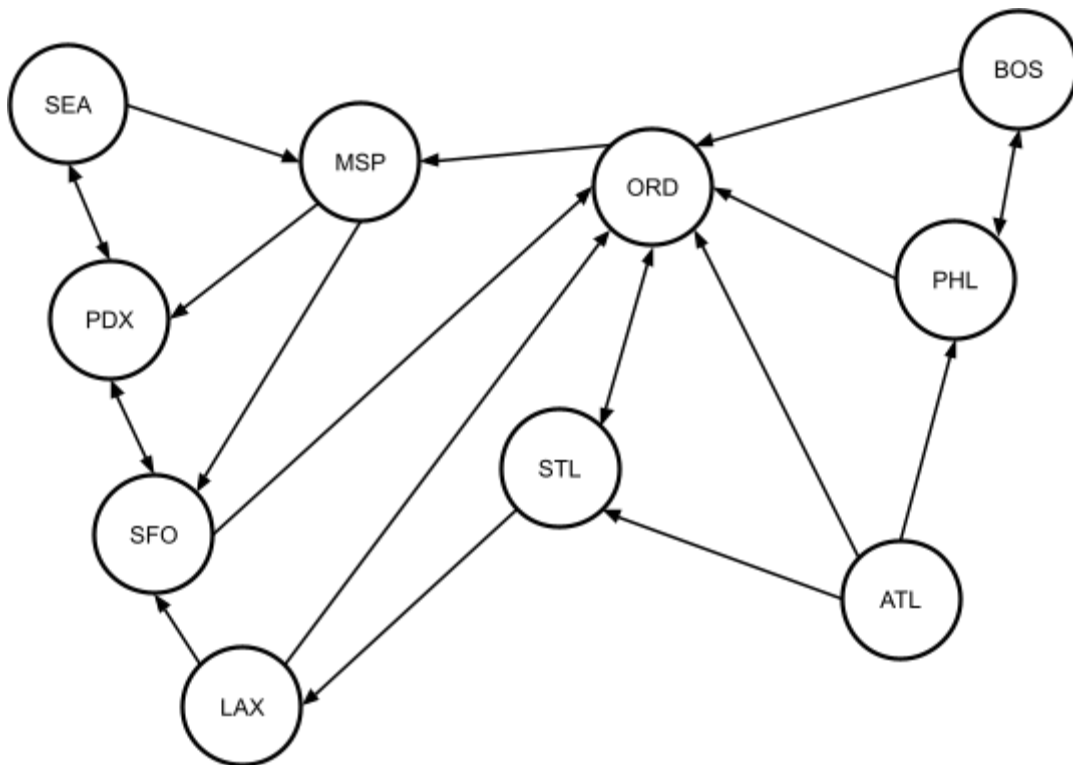
ATL: [ORD, PHL, STL],
BOS: [ORD, PHL],
LAX: [ORD, SFO, STL],
MSP: [ORD, PDX, SEA, SFO],
ORD: [ATL, BOS, LAX, MSP, PHL, SFO, STL],
PDX: [MSP, SEA, SFO],
PHL: [ATL, BOS, ORD],
SEA: [MSP, PDX],
SFO: [LAX, MSP, ORD, PDX],
STL: [ATL, LAX, ORD]

```

- As an adjacency matrix, the graph would look like this:

	ATL	BOS	LAX	MSP	ORD	PDX	PHL	SEA	SFO	STL
ATL	0	0	0	0	1	0	1	0	0	1
BOS	0	0	0	0	1	0	1	0	0	0
LAX	0	0	0	0	1	0	0	0	1	1
MSP	0	0	0	0	1	1	0	1	1	0
ORD	1	1	1	1	0	0	1	0	1	1
PDX	0	0	0	1	0	0	0	1	1	0
PHL	1	1	0	0	1	0	0	0	0	0
SEA	0	0	0	1	0	1	0	0	0	0
SFO	0	0	1	1	1	1	0	0	0	0
STL	1	0	1	0	1	0	0	0	0	0

- Note that this matrix is symmetric.
- What is the space complexity of each of these representations?
 - Adjacency list: $O(|V| + |E|)$
 - Adjacency matrix: $O(|V|^2)$
- Thus, the adjacency list is more space efficient when the graph is ***sparse***, i.e. when it has relatively few edges.
- What if our graph is a directed graph, e.g. if we have a flight from airport A to airport B but not a return flight?
- Each of these representations can still be used. For example, say we have this graph:



- Now, our adjacency list would look like this:

ATL: [ORD, PHL, STL],

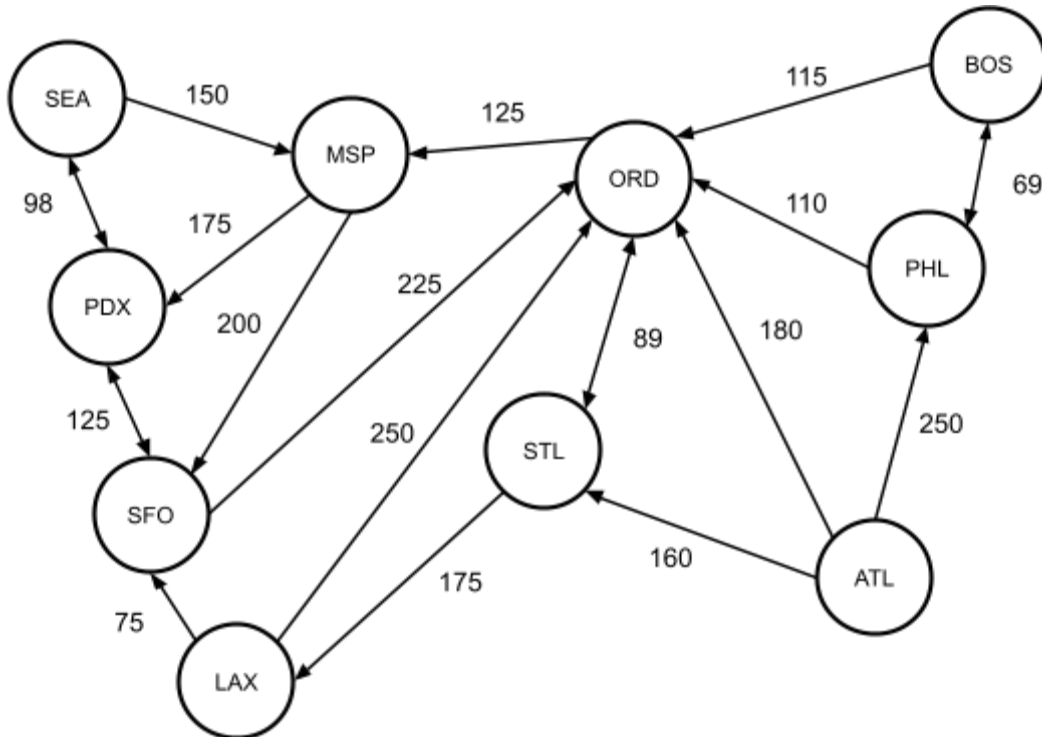
BOS: [ORD, PHL],
 LAX: [ORD, SFO],
 MSP: [PDX, SFO],
 ORD: [MSP, STL],
 PDX: [SEA, SFO],
 PHL: [BOS, ORD],
 SEA: [MSP, PDX],
 SFO: [ORD, PDX],
 STL: [LAX, ORD]

- Here, each vertex lists only the edges going *from* it.
- The adjacency matrix for this graph would look like this:

	ATL	BOS	LAX	MSP	ORD	PDX	PHL	SEA	SFO	STL
ATL	0	0	0	0	1	0	1	0	0	1
BOS	0	0	0	0	1	0	1	0	0	0
LAX	0	0	0	0	1	0	0	0	1	0
MSP	0	0	0	0	0	1	0	0	1	0
ORD	0	0	0	1	0	0	0	0	0	1
PDX	0	0	0	0	0	0	0	1	1	0
PHL	0	1	0	0	1	0	0	0	0	0
SEA	0	0	0	1	0	1	0	0	0	0
SFO	0	0	0	0	1	1	0	0	0	0
STL	0	0	1	0	1	0	0	0	0	0

- Note that this matrix is no longer symmetric.

- We can go one step further with each representation, incorporating weights. Say our graph contains the costs of flights between cities:



- Now, our adjacency list would store the weights/costs along with the edges:

```

ATL: [{ORD: 180}, {PHL: 250}, {STL: 160}],
BOS: [{ORD: 115}, {PHL: 69}],
LAX: [{ORD: 250}, {SFO: 75}],
MSP: [{PDX: 175}, {SFO: 200}],
ORD: [{MSP: 125}, {STL: 89}],
PDX: [{SEA: 98}, {SFO: 125}],
PHL: [{BOS: 69}, {ORD: 110}],
SEA: [{MSP: 150}, {PDX: 98}],
SFO: [{ORD: 225}, {PDX: 125}],
STL: [{LAX: 175}, {ORD: 89}]

```


- The adjacency matrix for this graph would now hold these weights/costs instead of just binary values:

	ATL	BOS	LAX	MSP	ORD	PDX	PHL	SEA	SFO	STL
ATL	0	0	0	0	180	0	250	0	0	160
BOS	0	0	0	0	115	0	69	0	0	0
LAX	0	0	0	0	250	0	0	0	75	0
MSP	0	0	0	0	0	175	0	0	200	0
ORD	0	0	0	125	0	0	0	0	0	89
PDX	0	0	0	0	0	0	0	98	125	0
PHL	0	69	0	0	110	0	0	0	0	0
SEA	0	0	0	150	0	98	0	0	0	0
SFO	0	0	0	0	225	125	0	0	0	0
STL	0	0	175	0	89	0	0	0	0	0

- We could also use a special value here (e.g. -1) to indicate there is no edge.

Single source reachability

- One important question we want to be able to ask about a graph is what nodes are reachable from some specific node.
- For example, we might ask about our graph above: what airports are reachable from PDX?
- We can use a very simple algorithm to answer this question. It looks like this, if we're trying to find reachable vertices from some vertex v_i :

1. Initialize an empty set of reachable vertices.
 2. Initialize an empty stack. Add v_i to the stack.
 3. If the stack is not empty, pop a vertex v from the stack.
 4. If v is not in the set of reachable vertices:
 - Add it to the set of reachable vertices.
 - Add each vertex that is direct successor of v to the stack.
 5. Repeat from 3.
- Looking for airports reachable from PDX would look like this:
 1. reachable: {}
stack: [PDX]
 2. v: PDX
successors: [SEA, SFO]
reachable: {PDX}
stack: [SEA, SFO]
 3. v: SFO
successors: [ORD, PDX]
reachable: {PDX, SFO}
stack: [SEA, ORD, PDX]
 4. v: PDX (already reachable)
successors: --
reachable: {PDX, SFO}
stack: [SEA, ORD]
 5. v: ORD
successors: [MSP, STL]
reachable: {ORD, PDX, SFO}
stack: [SEA, MSP, STL]

6. v: STL
successors: [LAX, ORD]
reachable: {ORD, PDX, SFO, STL}
stack: [SEA, MSP, LAX, ORD]
7. v: ORD (already reachable)
successors: --
reachable: {ORD, PDX, SFO, STL}
stack: [SEA, MSP, LAX]
8. v: LAX
successors: [ORD, SFO]
reachable: {LAX, ORD, PDX, SFO, STL}
stack: [SEA, MSP, ORD, SFO]
9. v: SFO, ORD (both already reachable)
successors: --
reachable: {LAX, ORD, PDX, SFO, STL}
stack: [SEA, MSP]
10. v: MSP
successors: [PDX, SFO]
reachable: {LAX, MSP, ORD, PDX, SFO, STL}
stack: [SEA, PDX, SFO]
11. v: SFO, PDX (both already reachable)
successors: --
reachable: {LAX, MSP, ORD, PDX, SFO, STL}
stack: [SEA]

```

12. v: SEA
    successors: MSP, PDX
    reachable: {LAX, MSP, ORD, PDX, SEA, SFO, STL}
    stack: [MSP, PDX]

13. v: PDX, MSP (both already reachable)
    Successors: --
    reachable: {LAX, MSP, ORD, PDX, SEA, SFO, STL}
    stack: []

14. Done (stack empty)
    reachable: {LAX, MSP, ORD, PDX, SEA, SFO, STL}

```

- This algorithm can be implemented using either the adjacency list representation or the adjacency matrix representation.
- We could also use a queue instead of a stack. This would result in a different order of exploration of the graph. We'll look at this in the next section.

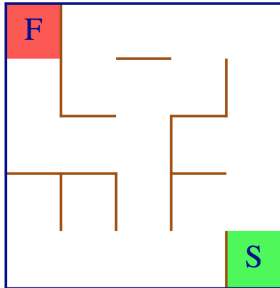
Depth-first search and Breadth-first search

- The reachability algorithm we saw above was an instance of ***depth-first search*** (or ***DFS***).
- DFS is an algorithm for exploring a tree where we travel a particular path as far as we can take it before trying another path.
 - In other words, in DFS, the neighbors of a node's neighbor are explored before exploring the node's other neighbors.
- DFS can be implemented using a stack, like the reachability algorithm above.

- If we replace the stack with a queue, that results in an exploration known as **breadth-first search** (or **BFS**).
- BFS explores a tree by traveling all paths to a given depth, then travelling all those paths one step deeper, then travelling them one step deeper, etc.
 - In other words, in BFS, all of a node's neighbors are explored before exploring its neighbors' neighbors.
 - That means BFS travels all paths of length 1, then travels all paths of length 2, then travels all paths of length 3, etc.
- The general algorithm for DFS and BFS is below. For DFS, we use a stack, and for BFS, we use a queue:
 1. Initialize an empty set of visited vertices.
 2. Initialize an empty stack (DFS) or queue (BFS). Add v_i to the stack/queue.
 3. If the stack/queue is not empty, pop/dequeue a vertex v .
 4. Perform any desired processing on v .
 - E.g. check if v meets a desired condition.
 5. (DFS only): If v is not in the set of visited vertices:
 - Add v to the set of visited vertices.
 - Push each vertex that is direct successor of v to the stack.
 6. (BFS only):
 - Add v to the set of visited vertices.
 - For each direct successor v' of v :
 - If v' is not in the set of visited vertices, enqueue it into the queue
 7. Repeat from 3.
- Often, we use BFS or DFS when we are looking for a node with a particular characteristic.

- For example, both algorithms can be used to find a path from start to finish in a maze.
- For example, this presentation shows how both DFS and BFS can be used to find the end of the maze below:

http://web.engr.oregonstate.edu/~hessro/static/media/GraphAlgorithm sII_DFS_BFS.d8183866.pdf



- We can make some comparisons between DFS and BFS:
 - DFS is a **backtracking** search: if we're looking for a node with a specific characteristic and DFS takes a path that doesn't contain such a node, it will backtrack to try a different path.
 - In an infinite graph, DFS can become lost down an infinite path without ever finding a solution.
 - BFS is **complete** and **optimal**: if a solution exists in the graph, BFS is guaranteed to find it, and it will find the shortest path to that solution.
 - However, BFS may take a long time to find a solution if the solution is deep in the graph.
 - DFS may find a deep solution more quickly.
 - Both algorithms have $O(V)$ space complexity in the worst case.
 - However, BFS may take up more space in practice.
 - If the graph has a high **branching factor**, i.e. if each node has many neighbors, BFS can take a lot of memory to maintain all of the paths it's exploring on the queue.

Dijkstra's algorithm: single source lowest-cost paths

- **Dijkstra's algorithm** finds the shortest/lowest-cost path from a specified vertex in a graph to all other reachable vertices in the graph.
- For example, in our graph above, you could use Dijkstra's algorithm to say not only *what* airports could be reached from PDX but also *what the cheapest cost* to reach each of those airports was.
- Dijkstra's algorithm is structured very much like DFS and BFS, except for this algorithm we will use a *priority queue* to order our search.
 - The priority values used in the queue correspond to the *cumulative* distance to each vertex added to the PQ.
 - Thus, we are always exploring the remaining node with the minimum cumulative cost.
- Here's the algorithm, which begins with some source vertex v_s :
 - Initialize an empty map/hash table representing visited vertices.
 - Key is the vertex v .
 - Value is the min distance d to vertex v .
 - Initialize an empty priority queue, and insert v_s into it with distance (priority) 0.
 - While the priority queue is not empty:
 - Remove the first element (a vertex) from the priority queue and assign it to v . Let d be v 's distance (priority).
 - If v is not in the map of visited vertices:
 - Add v to the visited map with distance/cost d .
 - For each direct successor v_i of v :
 - Let d_i equal the cost/distance associated with edge (v, v_i) .
 - Insert v_i to the priority queue with distance (priority) $d + d_i$.

- This version of the algorithm only keeps track of the minimum distance to each vertex, but it can be easily modified to keep track of the min-distance path, too.
 - Augment the visited vertex map and the priority queue to keep track of the vertex *previous* to each one added.
- The complexity of this version of the algorithm is $O(|E| \log |E|)$.
 - The innermost loop is executed at most $|E|$ times, and the cost of the instructions inside the loop is $O(\log |E|)$.
 - Inner cost comes from inserting into the PQ.